

Lab 7 - Project Report

AI24BTECH11031 - Shivram S
AI23BTECH11022 - Dhadheechi Ravva

Contents

1	Introduction	2
2	Usage	2
3	Modifications to the Simulator	2
4	Design of The Cache	2
5	Command Line Interface	4
6	Testing	4
7	Future Scope	5
8	File Structure	5

1 Introduction

As part of the CS2323 Computer Architecture course, we had to extend the RISC-V simulator developed in Lab 4 to support cache simulation. This report outlines the design of the cache simulator.

2 Usage

The project includes a Makefile which can be used to build the project and test it. To build, run:

```
$ make
```

This produces an executable called `riscv_sim` in the project directory.

The simulator can be run using the following command:

```
$ ./riscv_sim
```

3 Modifications to the Simulator

The `Simulator` struct had to be extended to support the cache simulator. Two fields - `cache` and `cache_config` were added to hold the cache and its configuration.

The `cache_enabled` field indicates whether the cache simulator is enabled. The `execution_in_progress` field indicates if the simulator is currently executing a program - this is necessary to ensure that the cache configuration is not changed during execution.

The implementation of load and store instructions were changed to use two new functions - `mem_read` and `mem_write`. These functions delegate memory access to the cache if it is enabled, otherwise they directly access the memory.

4 Design of The Cache

The simulator's cache consists of several **lines**. Each line consists of several **entries**, and each entry holds a block of memory.

The cache has a **replacement policy** which determines the block to be replaced in case of a conflict, and a **write-back policy**, which tells it how to handle writes.

The Cache struct holds all data related to the simulator's cache

```
typedef struct Cache {
    // Cache configuration
    size_t num_lines, block_size, associativity;
    enum {WRITEBACK, WRITETHROUGH} write_policy;
    enum {FIFO, LRU, RANDOM} replacement_policy;

    size_t hits, misses, writebacks; // Statistics

    size_t monotime; // For maintaining insert and access times
    uint8_t *mem; // Simulator memory
    CacheLine *lines; // The actual cache data
    FILE *output_file; // File for logging accesses
} Cache
```

The number of lines in the cache is given by `num_lines`, the number of entries in a line is given by `associativity`, and the number of bytes in a block is given by `block_size`.

Each cache entry has the format

```
typedef struct CacheEntry {
    int valid, dirty;
    uint64_t tag;
    union { uint64_t insert_time, access_time; };
    uint8_t *data;
} CacheEntry;
```

The `valid` bit determines if the entry is valid, and the `dirty` bit indicates that the entry contains changes not written to memory. The `insert_time` and `access_time` are used to select the entry to be replaced in the case of FIFO and LRU replacement policies respectively. The `tag` field holds the higher bits of the block's starting address.

The cache configuration is represented by a struct called `CacheConfig`, and is loaded from a file. The loading of the configuration is handled by the `load_cache_config` function.

```
typedef struct CacheConfig {
    size_t size, block_size, associativity,
    writeback_policy, replacement_policy;
} CacheConfig;
```

The logic of the cache itself is implemented in three functions.

- `cache_read` reads a specified number of bytes at an address.
- `cache_write` writes a specified number of bytes to an address.
- `cache_evict` is an internal function that selects blocks to be evicted from the cache.

In addition to this, there are several smaller functions for initializing the cache, to dump the cache to a file, and to invalidate the cache. The cache simulator also logs all memory access to a file.

5 Command Line Interface

All simulator commands are supported. The following cache-specific commands have been added:

- `cache_sim enable <config_file>`: Enables the cache simulator, using the configuration in `output_file`.
- `cache_sim disable`: Disables the cache simulator.
- `cache_sim status`: Displays the status of the cache, along with the cache configuration (if cache is enabled).
- `cache_sim invalidate`: Invalidates all cache entries
- `cache_sim dump <output_file>` Writes valid cache entries to `output_file`.
- `cache_sim stats`: Displays cache statistics such as hit count, and hit rate.

6 Testing

The project contains several test cases in the `test/` directory. The shell script `test.sh` runs each test case and compares the cache output to the expected output in `expected.output`. In addition to this, the number of hits, misses, and writebacks were compared manually.

7 Future Scope

- Implementing a real-time graphical view of the cache to facilitate profiling of code.
- Adding ability to view the content of specific blocks in the cache.
- Adding more replacement policies.
- Showing the 'dirty' status separately for each word in a block.

8 File Structure

```
+-- Makefile
+-- report.pdf
| \-- main.tex // Source file for the report
+-- src
| +-- asm // Source code for the assembler
| | +-- emitter.c
| | +-- emitter.h
| | +-- lexer.c
| | +-- lexer.h
| | +-- parser.c
| | +-- parser.h
| | +-- tables.c
| | \-- tables.h
| +-- main.c
| +-- simulator.c // Source code for the simulator
| +-- simulator.h
| +-- cache.c // Source code for the cache simulator
| +-- cache.h
+-- test // Testcases
\-- test.sh // Automatic testing script
```